



HALMSTAD UNIVERSITY

BIG DATA PARALLEL PROGRAMMING

Project Report

Shankaranarayanan Bangalore Ramalingam

Table of Contents

Problem Statement:	3
Introduction:	3
Discussion.....	3
Apache Spark:	4
Watson Studio:.....	4
Creating a project in IBM Watson Studio	5
Uploading Dataset in IBM Watson.....	5
Creating Jupyter Notebook	6
Pyspark.....	6
Exploratory Analysis.....	7
Showing Data	8
Handling missing data and upsampling	9
Handling Missing Data	10
Data Filtering.....	11
Graphical Representation	12
String Indexer, OneHotEncoderEstimator	13
Vector Assembler. Standard Scalar.....	14
Splitting Data.....	14
Machine Learning.....	14
Logistic Regression	14
Decision Tree.....	15
Random Forest Tree.....	16
Model Selection:	16
Creating Output Link on Jupyter Notebook	16
Running as Job in Watson Studio.....	17
Comparing local performance and cloud performance.....	17

Problem Statement:

The aim of this project is predict if a person will be getting an ischemic heart disease and stroke or not, and this is based on several attributes from the previous medical data. This data was taken from a GitHub page, but this was previously even a kaggle competition.

Keywords: Class imbalance, IBM Watson studio, Decision Tree, Logistic Regression, Random Forrest

Introduction:

As I had mentioned in the Problem statement section, we are going to predict probability of a person facing a heart stroke. According to the world health organization, ischemic heart disease and stroke is one of the largest killers in the world. This is caused by different attributes such as places they live, smoking habits, type of work etc. We will try to predict this based on some attributes provided in the dataset.

In this project we are using supervised machine learning classification algorithms to solve this. In supervised machine learning each data attribute is a labelled class and we have to train a model by using different algorithms and predict the data that may have lack of class label. In this heart stroke data we can see that there are 43400 entries and has 12 attributes totally. Each entry in the dataset consist of information about an individual. In order to perform this task of classification I have done the task using pyspark to understand the dataset and also implement the machine learning algorithm. The cloud platform used in the project is the IBM Watson studio.

I had created a project in the IBM Watson studio, following which I had selected an environment in which python and spark is installed already. After created the environment, I had uploaded the input file (train and test data) within the project. I had added a new jupyter notebook in the cloud and loaded the data in the jupyter notebook. Following which I had added performed explanatory analysis in which there was the process of cleaning data, handling class imbalance. After explanatory analysis there are many string attributes in the dataset which must be converted to numerical attributes. In order to perform this, I have used the StringIndexer. After this I have used the vector assembler and standard scaler in the dataset.

I have implemented 3 machine learning algorithms in the dataset, namely they are Logistic Regression, Decision Tree algorithm and Random forest algorithm. Each of the algorithm will return a ROC value, following which I have applied hyper parameter tuning on each of the algorithm to increase the ROC value by choosing the best parameter.

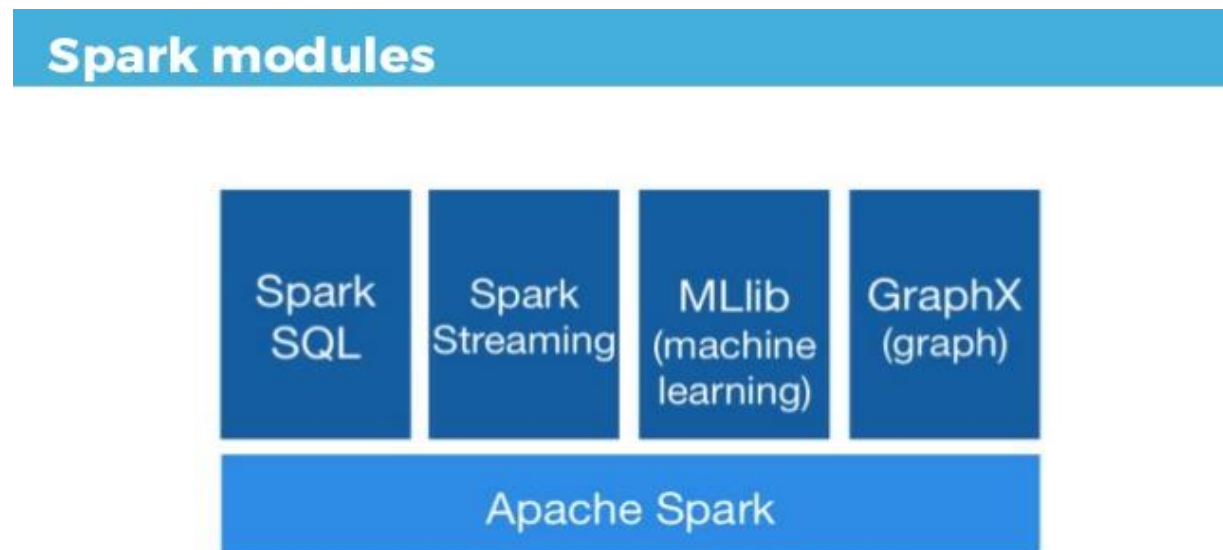
Following which I had run the whole setup as a job in IBM Watson studio, and logs can also be monitored in the cloud environment itself. One of the main reasons for choosing the IBM Watson studio is the ease of use, and presence of a lot of data connectors.

Discussion

In this part the whole project methodology will be explained that means the flow of project including IBM Watson studio configuration, loading dataset, explanatory analysis, data cleaning, handling class imbalance, upsampling, StringIndexer, Vector Assembler, standard scaler, machine learning model implementation (3 namely). All the algorithms are evaluated based on ROC and then we select the best model to make the prediction on this basis. After which we will apply each of the models are run on the test data as well.

Apache Spark:

Apache spark is the open source project which provides cluster computing framework. Spark provides the interface for programming entire clusters with implicit data parallelism and fault tolerance. It is used for large scale data analytics as provides features such as parallelism. Apache spark can also be used to implement machine learning algorithm with the help of the MLlib library (This is part of the pyspark library of Python). Apache Spark has many modules which can be used to implement the machine learning models on the data. The screenshot of the spark modules is given below



But here in this project we are only using the Spark SQL and MLlib module in spark as mentioned above. Hence, we will only provide explanation for those modules.

Spark SQL: This module was built on top of the spark core, and it provides a data abstraction called Dataframes, which provides structured and semi-structured data. The Spark SQL module also provides the SQL language support, with command-line interfaces and JDBC/ODBC servers.

MLlib: Apache Spark provides a Machine learning API called MLlib. Pyspark has this machine learning API in Python as well. Many common machine learning and statistical algorithms have been implemented and shipped with MLlib which simplifies large scale machine learning pipelines, which includes classification, regression, feature extraction, transformation etc.

Watson Studio:

The Watson studio is a cloud platform which helps data Scientists to prepare data and build models in the cloud which can be scaled. The Watson studio provides the suits which are necessary to accelerate the machine learning and deep learning workflows, which can used to infuse AI into the business development model.

Creating a project in IBM Watson Studio

First creating an account in the Watson studio, we can create a project with the help of the “lite” plan. As this is the free configuration. I had created a project with name of “BDPP_Project”.

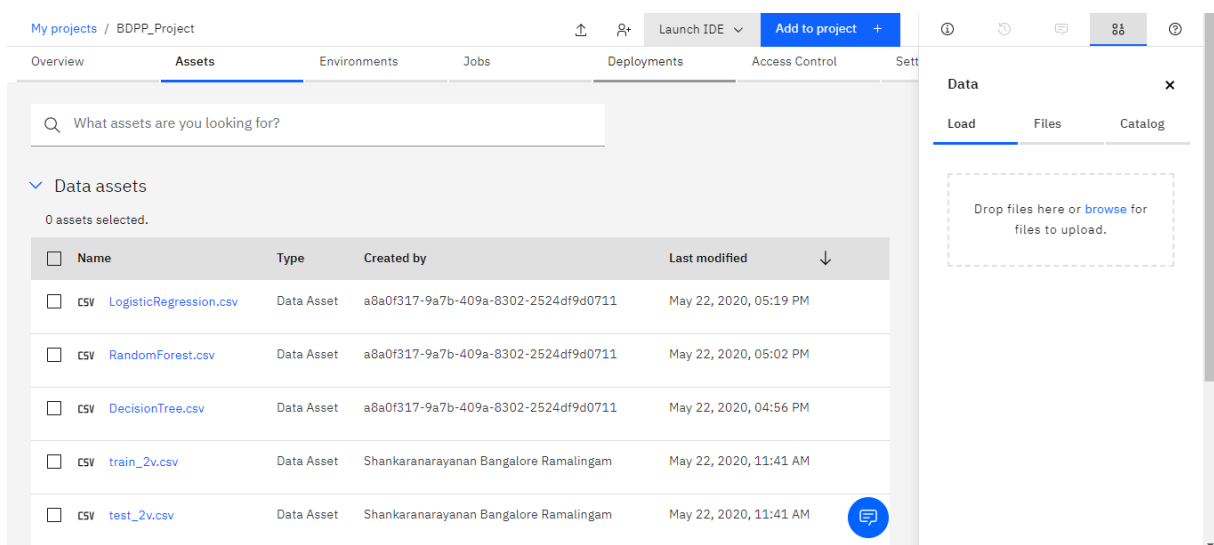
Name	Role	Collaborators	Date created	Last updated
BDPP_Project	Admin		May 22, 2020	May 22, 2020

This project has an environment of Python 3.6. & Apache Spark 2.4, and an automatic cloud storage was created by the IBM Watson studio. Attaching the configuration screenshot below

Environment	Default Spark 2.4 & Python 3.6
Creator	IBM
Hardware configuration (Driver)	1 vCPU and 4 GB RAM
Hardware configuration (Executor)	1 vCPU and 4 GB RAM
Number of executors	2
Spark version	2.4
Software version	Python 3.6

Uploading Dataset in IBM Watson

After creating an account and selecting the environment, the input files should be uploaded in the cloud, in Watson studio, there will be assets tab created after project creation. The input files have to be uploaded into here by clicking in the “New data asset” button. This is open a menu where we can upload the data into the cloud.



The screenshot shows the IBM Watson Studio interface. The top navigation bar includes 'My projects / BDPP_Project', 'Launch IDE', and 'Add to project'. The main tabs are Overview, Assets (selected), Environments, Jobs, Deployments, Access Control, and Settings. The Assets tab displays a search bar and a list of data assets. A right-hand panel shows a 'Data' dialog with a 'Load' tab and a file upload area.

Name	Type	Created by	Last modified
LogisticRegression.csv	Data Asset	a8a0f317-9a7b-409a-8302-2524df9d0711	May 22, 2020, 05:19 PM
RandomForest.csv	Data Asset	a8a0f317-9a7b-409a-8302-2524df9d0711	May 22, 2020, 05:02 PM
DecisionTree.csv	Data Asset	a8a0f317-9a7b-409a-8302-2524df9d0711	May 22, 2020, 04:56 PM
train_2v.csv	Data Asset	Shankaranarayanan Bangalore Ramalingam	May 22, 2020, 11:41 AM
test_2v.csv	Data Asset	Shankaranarayanan Bangalore Ramalingam	May 22, 2020, 11:41 AM

The new file is uploaded by clicking on the browse button, and we can upload the data into the cloud. In the above screenshot we can see that the input files are loaded in the tab, and this will serve as the bucket for our project.

Creating Jupyter Notebook

After uploading all the input files into the cloud, we have to create a notebook in which we will writing code, and where all the ML algorithms are implemented. In order to create a jupyter notebook in IBM Watson studio, click on the new Notebook button in the assets tab. Following which name of the Jupyter notebook is provided, then the runtime environment is chosen and create button is clicked. Attaching the screenshot below

New notebook

Blank From file From URL

Name
Type notebook name here

Description (optional)
Type your description here

Select runtime
Default Spark 2.4 & Python 3.6 (Driver: 1 vCPU 4 GB RAM, 2 Exe

The selected runtime uses 1 driver with 1 vCPU and 4 GB RAM, and 2 executors each with 1 vCPU and 4 GB RAM. It consumes 1.5 capacity units per hour. [Learn more](#) about capacity unit hours and Watson Studio pricing plans.

Language
☒ Python 3.6

Spark version
☒ 2.4

Cancel Create

Pyspark

After creating the notebook in the cloud, I had loaded the data into the Jupyter created in the cloud, attaching the screenshot of loading data in the cloud environment

```
In [193]: import ibmos2spark
# @hidden_cell
credentials = {
    'endpoint': 'https://s3-api.us-geo.objectstorage.service.networklayer.com',
    'service_id': 'iam-ServiceId-aaf06632-a925-4f1b-a1ad-4eb3f5182d4c',
    'iam_service_endpoint': 'https://iam.cloud.ibm.com/oidc/token',
    'api_key': 'cfdE0b2otcGpFI7mW6b1nD_tZKjR9z-zqgR0rxPmu2a'
}

configuration_name = 'os_4fe49bf4d14c41f9a72c9ef73f278c1a_configs'
cos = ibmos2spark.CloudObjectStorage(sc, credentials, configuration_name, 'bluemix_cos')

from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
heart_stroke_train = spark.read\
    .format('org.apache.spark.sql.execution.datasources.csv.CSVFileFormat')\
    .option('header', 'true')\
    .load(cos.url('train_2v.csv', 'bdppproject-donotdelete-pr-bibcnvcar6fhub'))
heart_stroke_train.take(5)
```

Here we are using the IBM cloud object storage feature in the Watson studio, which will help us load the csv in the jupyter notebook. Similarly, even the test data is also loaded in the cloud.

```
testSet = spark.read\
    .format('org.apache.spark.sql.execution.datasources.csv.CSVFileFormat')\
    .option('header', 'true')\
    .load(cos.url('test_2v.csv', 'bdppproject-donotdelete-pr-bibcnvcar6fhub'))
testSet.take(5)
```

Exploratory Analysis

After loading the data into the cloud, I am doing exploratory analysis. Firstly, I am checking the data length and checking the schema for both the train and test data.

```
heart_stroke_train.count(),len(heart_stroke_train.columns)
(43400, 12)
```

```
testSet.count(),len(testSet.columns)
(18601, 11)
```

```
heart_stroke_train.printSchema()
```

```
root
|-- id: string (nullable = true)
|-- gender: string (nullable = true)
|-- age: string (nullable = true)
|-- hypertension: string (nullable = true)
|-- heart_disease: string (nullable = true)
|-- ever_married: string (nullable = true)
|-- work_type: string (nullable = true)
|-- Residence_type: string (nullable = true)
|-- avg_glucose_level: string (nullable = true)
|-- bmi: string (nullable = true)
|-- smoking_status: string (nullable = true)
|-- stroke: string (nullable = true)
```

```
testSet.printSchema()
```

```
root
|-- id: string (nullable = true)
|-- gender: string (nullable = true)
|-- age: string (nullable = true)
|-- hypertension: string (nullable = true)
|-- heart_disease: string (nullable = true)
|-- ever_married: string (nullable = true)
|-- work_type: string (nullable = true)
|-- Residence_type: string (nullable = true)
|-- avg_glucose_level: string (nullable = true)
|-- bmi: string (nullable = true)
|-- smoking_status: string (nullable = true)
```

We can see that 12 and 11 attributes in the train and test set respectively. The data description is shown below in the table

Columns	Description
ID	Patient ID
Gender	Gender of the Patient
Age	Age of the Patient
Hypertension	1-Incase suffering from hypertension, 0-No hypertension
Heart Disease	1- In case of suffering from heart disease,0- No heart disease
Ever_married	Yes/No
Work_type	Type of Occupation
Residence Type	Area of residence(Rural, Urban)
Avg_glucose_level	Average_glucose_level(measured after meal)
BMI	Body mass index
Smoking_status	Patient's smoke status
Stroke	0-no stroke, 1-suffered Stroke

Showing Data

The following screenshot shows the top 20 rows in the train and test dataset

```
heart_stroke_train.show()
```

id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
30669	Male	3.0	0	0	No	children	Rural	95.12	18.0	null
30468	Male	58.0	1	0	Yes	Private	Urban	87.96	39.2	never smoked
16523	Female	8.0	0	0	No	Private	Urban	110.89	17.6	null
56543	Female	70.0	0	0	Yes	Private	Rural	69.04	35.9	formerly smoked
46136	Male	14.0	0	0	No	Never_worked	Rural	161.28	19.1	null
32257	Female	47.0	0	0	Yes	Private	Urban	210.95	50.1	null
52800	Female	52.0	0	0	Yes	Private	Urban	77.59	17.7	formerly smoked
41413	Female	75.0	0	1	Yes	Self-employed	Rural	243.53	27.0	never smoked
15266	Female	32.0	0	0	Yes	Private	Rural	77.67	32.3	smokes
28674	Female	74.0	1	0	Yes	Self-employed	Urban	205.84	54.6	never smoked
10460	Female	79.0	0	0	Yes	Govt_job	Urban	77.08	35.0	null
64908	Male	79.0	0	1	Yes	Private	Urban	57.08	22.0	formerly smoked
63884	Female	37.0	0	0	Yes	Private	Rural	162.96	39.4	never smoked
37893	Female	37.0	0	0	Yes	Private	Rural	73.5	26.1	formerly smoked
67855	Female	40.0	0	0	Yes	Private	Rural	95.04	42.4	never smoked
25774	Male	35.0	0	0	No	Private	Rural	85.37	33.0	never smoked
19584	Female	20.0	0	0	No	Private	Urban	84.62	19.7	smokes
24447	Female	42.0	0	0	Yes	Private	Rural	82.67	22.5	never smoked
49589	Female	44.0	0	0	Yes	Govt_job	Urban	57.33	24.6	smokes
17986	Female	79.0	0	1	Yes	Self-employed	Urban	67.84	25.2	smokes

only showing top 20 rows


```
testSet.show()
```

id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
36306	Male	80.0	0	0	Yes	Private	Urban	83.84	21.1	formerly smoked
61829	Female	74.0	0	1	Yes	Self-employed	Rural	179.5	26.0	formerly smoked
14152	Female	14.0	0	0	No	children	Rural	95.16	21.2	null
12997	Male	28.0	0	0	No	Private	Urban	94.76	23.4	null
40801	Female	63.0	0	0	Yes	Govt_job	Rural	83.57	27.6	never smoked
9348	Female	66.0	1	0	Yes	Private	Urban	219.98	32.2	never smoked
51550	Female	49.0	0	0	Yes	Self-employed	Rural	74.03	25.1	null
60512	Male	46.0	0	0	Yes	Govt_job	Urban	120.8	32.5	never smoked
31309	Female	75.0	0	0	Yes	Self-employed	Rural	78.71	28.0	never smoked
39199	Male	75.0	0	0	Yes	Self-employed	Urban	77.2	25.7	smokes
15160	Female	17.0	0	0	No	Private	Rural	78.16	21.9	null
21705	Female	10.0	0	0	No	children	Urban	107.23	19.4	null
19042	Female	47.0	0	0	Yes	Private	Rural	91.6	26.7	never smoked
12249	Female	42.0	0	0	Yes	Private	Urban	83.05	32.3	null
33104	Female	67.0	0	0	Yes	Govt_job	Urban	236.6	24.2	never smoked
55264	Female	52.0	0	0	No	Self-employed	Urban	109.49	24.5	never smoked
29445	Male	73.0	0	0	Yes	Self-employed	Rural	109.66	40.0	null
49013	Female	19.0	0	0	No	Private	Rural	88.51	22.1	null
276	Male	15.0	0	0	No	children	Rural	101.36	22.3	null
47721	Female	37.0	0	0	Yes	Govt_job	Urban	165.44	36.1	formerly smoked

only showing top 20 rows

Handling missing data and upsampling

In this dataset, there are two output classes 0 and 1, and in the data 1 output class has 763 entries while the 0 output class has 42617, attaching the screenshot below



Now in order to handle the imbalance in data we have performed upsampling of the 1 output class.

```
heart_stroke_train_pos=heart_stroke_train.filter(heart_stroke_train["label"]==1)
heart_stroke_train_neg=heart_stroke_train.filter(heart_stroke_train["label"]==0)
```

The dataset is split into 2 classes which is 1 and 0, but here we will only upsample the 1 class, to handle the imbalance,

```
from pyspark.sql.functions import col
heart_stroke_train_pos= heart_stroke_train_pos.where(col('label')==1).sample(True,(42617/763)*3.0, seed = 2018)
heart_stroke_train_neg=heart_stroke_train_neg.where(col('label')==0)
heart_stroke_train = heart_stroke_train_neg.union(heart_stroke_train_pos)
```


```
heart_stroke_train_neg=heart_stroke_train_neg.where(col('label')==0).sample(True, 3.0, seed = 2018)
heart_stroke_train = heart_stroke_train_neg.union(heart_stroke_train_pos)
```

As shown above the data is split into 2 classes and we have used the sample method to upsample the data. We can see that the fraction multiple is negative outcomes by positive outcomes.

But in order to scale up the data, we can see that the whole data is multiplied 3 times, hence the number of entries is increased in both the classes, and this will result in both class entries having almost equal number of entries, the below output justifies that

```
+-----+-----+
|label|count|
+-----+-----+
|    1|128078|
|    0|128061|
+-----+-----+
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fb0f58af978>



Handling Missing Data


In this dataset, the columns “smoking_status” and “bmi” have null values in them, these null values have to be removed from the dataset. But the strategy used in both the columns are different, this is due to the datatype of both the columns. From the above **printSchema** and **dtypes** image we can see that “smoking_status” column is of string type and “bmi” is of integer type.

For the “smoking_status” column we have used to the **fillna** method provided by the dataframe object in spark. This method will replace all the null string to “No Info”. The below figure shows the table before and after using fillna method.

```
heart_stroke_train.filter(heart_stroke_train['label']==1).groupBy('smoking_status').count().orderBy('count',ascending=False).show()
smoking_status_plot=heart_stroke_train.filter(heart_stroke_train['label']==1).groupBy('smoking_status').count().orderBy('count',ascending=False)
data=smoking_status_plot.toPandas()
data.plot(x='smoking_status',y='count',kind='barh')
```

```
+-----+-----+
| smoking_status|count|
+-----+-----+
| never smoked|46593|
| formerly smoked|35945|
| null|23628|
| smokes|21912|
+-----+-----+
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fb0f685d898>




As we can see that there are 23628 entries which are in the “null” value section, and this null has been replaced with “No Info”.

```
smoking_status_plot=updated_info.filter(updated_info['label']==1).groupBy('smoking_status').count().orderBy('count',ascending=False)
updated_info.filter(updated_info['label']==1).groupBy('smoking_status').count().orderBy('count',ascending=False).show()
data=smoking_status_plot.toPandas()
data.plot(x='smoking_status',y='count',kind='barh')
```

```
+-----+-----+
| smoking_status|count|
+-----+-----+
| never smoked|46593|
| formerly smoked|35945|
| No Info|23628|
| smokes|21912|
+-----+-----+
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fb130034f60>



After this I noticed that the BMI column has null values present in it, I had replaced the null values present in the bmi column with the mean value.

bmi	bmi
36.6	29.202138491982318
null	31.4
32.5	30.9
34.4	26.6
24.0	26.4
29.0	31.5
27.4	30.7
22.8	27.5
null	28.5
24.2	25.8
29.7	24.2
36.8	25.6
27.3	27.3
null	25.0
28.2	30.3
30.9	30.6
37.5	29.3
25.8	28.0
37.8	27.0
null	29.7

The above screenshots shows the null value and the null values replaced after using the mean method on the dataframe. The below screenshot shows the code which was used to replace the value.

```
from pyspark.sql.functions import mean
mean = updated_info.select(mean(updated_info['bmi'])).collect()
mean_bmi = mean[0][0]
updated_info = updated_info.fillna(mean_bmi,['bmi'])
updated_info.filter(updated_info['label']==1).select('bmi')
```

DataFrame[bmi: double]

```
testSet=testSet.na.fill(mean_bmi,['bmi'])
```

Data Filtering

Here you can see the different filter data and their graphical representations.

work_type	count
Private	72171
Self-employed	40901
Govt_job	14695
children	311

ever_married	count
Yes	115023
No	13055

Residence_type	count
Urban	65082
Rural	62996

age	count
82.0	5829
81.0	7166
80.0	7986
79.0	11350
78.0	9512
77.0	4018
76.0	3865
75.0	3715
74.0	3930
73.0	2418
72.0	3442
71.0	3167
70.0	4078
69.0	3266
68.0	3405
67.0	3753
66.0	2733
65.0	2989
64.0	1356
63.0	2947


only showing top 20 rows

bmi	count
29.202138491982318	22764
31.4	1468
30.9	1459
26.6	1438
26.4	1406
31.5	1328
30.7	1312
27.5	1276
28.5	1228
25.8	1160
24.2	1144
25.6	1016
27.3	1006
30.3	999
25.0	999
30.6	993
29.3	983
28.0	971
27.0	961
29.7	947


only showing top 20 rows

Graphical Representation


<matplotlib.axes._subplots.AxesSubplot at 0x11b1718d0>



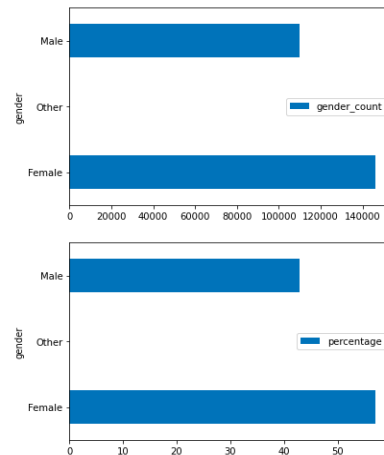
<matplotlib.axes._subplots.AxesSubplot at 0x11b863278>



<matplotlib.axes._subplots.AxesSubplot at 0x11b988908>



<matplotlib.axes._subplots.AxesSubplot at 0x11cacf390>



String Indexer, OneHotEncoderEstimator

The machine learning algorithms prefer to working with numbers than string values, and in our dataset we have a lot of string attributes, the StringIndexer will help us replace the string values with numbers, if we look at the our schema we can see that the following columns were of string attributes.

1. Gender
2. Ever_married
3. Work_type
4. Residence_type
5. Smoking_status

The OneHotEncoderEstimator will be useful to create continuous values and this will create categorical features. The below screenshot shows the implementation of the gender column being changed to numerical attributes.

```
#StringIndexer for gender column
from pyspark.ml.feature import StringIndexer
from pyspark.ml.feature import OneHotEncoderEstimator
gender_indexer=StringIndexer(inputCol="gender",outputCol="genderIndex")
updated_info=gender_indexer.fit(updated_info).transform(updated_info)
updated_info.select(["gender","genderIndex"]).show(5)
```

```
+-----+-----+
|gender|genderIndex|
+-----+-----+
| Male|          1.0|
| Male|          1.0|
| Male|          1.0|
|Female|          0.0|
|Female|          0.0|
+-----+-----+
only showing top 5 rows
```

```
#OneHotEncoderEstimate for the gender column
```

```
gender_encoder=OneHotEncoderEstimator(inputCols=["genderIndex"],outputCols=["genderVec"])
model=gender_encoder.fit(updated_info)
updated_info=model.transform(updated_info)
updated_info.select("genderVec").take(5)
```

```
[Row(genderVec=SparseVector(2, {1: 1.0})),
 Row(genderVec=SparseVector(2, {1: 1.0})),
 Row(genderVec=SparseVector(2, {1: 1.0})),
 Row(genderVec=SparseVector(2, {0: 1.0})),
 Row(genderVec=SparseVector(2, {0: 1.0}))]
```

The above 2 screenshots are only an example of one of the attributes, similarly the String Indexer and OneHotEncoderEstimator has been implemented for all the String attributes.

Vector Assembler. Standard Scaler

Vector assembler is used to combine all the attributes in one feature col and this feature or scaled features col is used for training the model as well.

Stand scaler is used for scale the values in feature col. If there are values that have much difference, then the stand scaler scale all the values and generate scaled features col. You can see the following snap.

features	scaledFeatures
(11, [1,2,6,7,9], [...]	[-1.1536667115755...
[0.0,1.0,58.0,1.0...	[-1.1536667115755...
[0.0,1.0,58.0,1.0...	[-1.1536667115755...
(11, [0,2,5,6,7,9]...	[0.86679808460811...
(11, [0,2,5,6,7,9]...	[0.86679808460811...

only showing top 5 rows

Splitting Data

Here I am splitting the training dataframe into train dataset and the validation set. In my project the data is split in the 70:30 ratio, wherein 70 is the train dataset and 30 is the validation set. After splitting the data, the number of entries in each of the sets are provided in the below screenshot

```

: val_data.count(), len(val_data.columns)
: (76773, 24)

: train_data.count(), len(train_data.columns)
: (179366, 24)

```

Machine Learning

Machine learning basically performs some rules on the dataset and make prediction of test data. This is supervised machine learning task and in supervised machine learning the data set always given with the label class. Now the data has been ready for the processing for machine leaning. Now I will use different number of algorithm and will find the ROC that will help to evaluate the algorithm.

Logistic Regression

The logistic Regression is classification algorithm, wherein we predict the label class by using the information provided in the dataset. I train the model in using the training data, and then run the prediction in the test set. But to find the ROC of the model, I have used the validation dataset which I had created. The results dataframe from the test set for each entry is shown below.

label	prediction	probability	age	work_type
0	0.0	[0.80379707391931...	37.0	Private
0	0.0	[0.96184693802282...	21.0	Private
0	1.0	[0.22446153607421...	79.0	Private
0	0.0	[0.78616625016903...	44.0	Private
0	0.0	[0.67685275119944...	54.0	Self-employed
0	0.0	[0.60146018483361...	56.0	Private
0	1.0	[0.41770503525074...	67.0	Private
0	0.0	[0.98620092203542...	5.0	children
0	0.0	[0.98620092203542...	5.0	children
0	1.0	[0.45922127555439...	68.0	Private
0	1.0	[0.45922127555439...	68.0	Private
0	0.0	[0.98675212113989...	5.0	children
0	0.0	[0.97903799113763...	13.0	Private
0	0.0	[0.97903799113763...	13.0	Private
0	0.0	[0.94493038029532...	26.0	Private
0	0.0	[0.94493038029532...	26.0	Private
0	1.0	[0.33123146765210...	52.0	Self-employed
0	0.0	[0.90034772131555...	36.0	Private
0	1.0	[0.20142166051380...	74.0	Govt_job
0	0.0	[0.79407088763163...	46.0	Govt_job

Decision Tree

Decision tree algorithm belongs to a family of supervised machine learning algorithm. The main goal of the decision tree algorithm is predict the output (class or target value) by learning simple decision rules inferred from prior data. Like Logistic Regression we have used train dataset to train model, and done the prediction in the test model. The ROC of the model was 0.82 before hyperparameter tuning, and after that ROC was 0.88

DataFrame[label: int, prediction: double, probability: vector,

label	prediction	probability	age	work_type
0	1.0	[0.15041696695150...	80.0	Self-employed
0	1.0	[0.15041696695150...	80.0	Self-employed
0	0.0	[0.80543530543530...	37.0	Private
0	0.0	[1.0,0.0]	21.0	Private
0	0.0	[1.0,0.0]	44.0	Private
0	1.0	[0.27169359664871...	79.0	Private
0	1.0	[0.27169359664871...	79.0	Private
0	1.0	[0.48888888888888...	44.0	Private
0	1.0	[0.48888888888888...	44.0	Private
0	0.0	[1.0,0.0]	34.0	Private
0	1.0	[0.15882967607105...	54.0	Self-employed
0	0.0	[0.60396039603960...	67.0	Private
0	0.0	[1.0,0.0]	5.0	children
0	0.0	[1.0,0.0]	5.0	children
0	0.0	[1.0,0.0]	5.0	children
0	0.0	[0.60396039603960...	68.0	Private
0	0.0	[1.0,0.0]	4.0	children
0	0.0	[1.0,0.0]	4.0	children
0	1.0	[0.48888888888888...	44.0	Private
0	0.0	[1.0,0.0]	5.0	children

only showing top 20 rows

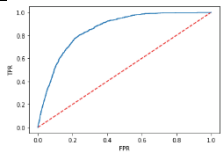
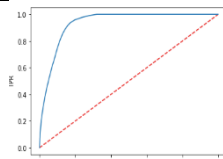
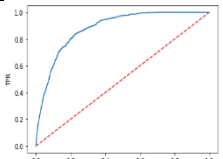
Random Forest Tree

Random forest, like its name implies, consists of many individual decision trees that operate as an ensemble. Each individual tree in the random forest spits out a class prediction and the class with the most votes becomes our model's prediction. The ROC for this 0.85 and after parameter tuning it is 0.88.

label	prediction	probability	age	work_type
0	0.0	[0.85518521905847...	37.0	Private
0	0.0	[0.86267010483866...	21.0	Private
0	1.0	[0.24102076143563...	79.0	Private
0	0.0	[0.81515478841869...	44.0	Private
0	0.0	[0.69635537790376...	54.0	Self-employed
0	0.0	[0.61814925040722...	56.0	Private
0	1.0	[0.29385888361690...	67.0	Private
0	0.0	[0.90740495233839...	5.0	children
0	0.0	[0.90740495233839...	5.0	children
0	1.0	[0.30806297945393...	68.0	Private
0	1.0	[0.30806297945393...	68.0	Private
0	0.0	[0.94025611990864...	5.0	children
0	0.0	[0.91534075092428...	13.0	Private
0	0.0	[0.91534075092428...	13.0	Private
0	0.0	[0.90205032490179...	26.0	Private
0	0.0	[0.90205032490179...	26.0	Private
0	1.0	[0.44685142955909...	52.0	Self-employed
0	0.0	[0.83269488014656...	36.0	Private
0	1.0	[0.19920038422445...	74.0	Govt_job
0	0.0	[0.81672608781590...	46.0	Govt_job

Model Selection:

We can see that, there were 3 models which were implemented, and the best models are chosen based on the ROC of the each of the models. The below table compares different models which are used in the project.

Category	Logistic Regression	Decision Tree Classification	Random Forest Classification
AreaUnderROC before Parameter tuning	0.8515995624882962	0.8601771331423859	0.8672775922191699
AreaUnderROC after Parameter tuning	0.8519365138452384	0.8901265969186091	0.880087906275998
ROC Curve after hyperparameter tuning			

Creating Output Link on Jupyter Notebook

Finally, after choosing best model, and the prediction being applied for each of the models in the test data. After that the predicted values from the test dataframe will be stored in csv file in the cloud assests.


The below screenshot shows the output csv stored in the cloud, and they can be downloaded from there as csv.

<input type="checkbox"/>	Name	Type	Created by	Last modified	↓
<input type="checkbox"/>	CSV LogisticRegression.csv	Data Asset	a8a0f317-9a7b-409a-8302-2524df9d0711	May 22, 2020, 08:49 PM	
<input type="checkbox"/>	CSV RandomForest.csv	Data Asset	a8a0f317-9a7b-409a-8302-2524df9d0711	May 22, 2020, 08:32 PM	
<input type="checkbox"/>	CSV DecisionTree.csv	Data Asset	a8a0f317-9a7b-409a-8302-2524df9d0711	May 22, 2020, 08:26 PM	

We can see that this is done for each of the models.

Running as Job in Watson Studio

After completing the code in the jupyter notebook, i have submitted the cluster job with main python file which is present in the cluster, we can see that logs of the job submitted in the cloud, the screenshot containing the job details is shown below.

Job name	Associated asset	Last run	Started by	Created by
This is the result	Notebook	 Finished Jun 13, 2020, 01:01 AM	Shankaranarayanan Bangalore Ramalingam Jun 13, 2020, 12:40 AM	Shankaranarayanan Bangalore Ramalingam

The below screenshot shows the logs generated when running the job in the cloud.

...	
[I 2020-06-12 19:04:18.243 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type status on channel iopub
[I 2020-06-12 19:04:18.249 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	Sending Shell Message type execute_request
[I 2020-06-12 19:04:18.324 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type status on channel iopub
[I 2020-06-12 19:04:18.335 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type execute_input on channel iopub
[I 2020-06-12 19:04:20.863 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:25.867 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:33.867 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:34.875 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:38.878 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:43.898 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:44.892 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:45.898 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:46.101 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:47.888 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:48.182 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:49.186 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:50.186 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:51.189 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:52.118 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:53.119 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:54.113 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:55.899 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:56.118 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:57.109 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub
[I 2020-06-12 19:04:58.122 NotebookApp] [b179f840-e4e8-45df-a106-ade5693122f0]	registered message type spark_monitor_msg on channel iopub

Comparing local performance and cloud performance

The main aim of using the project is cloud is to have a better performance, where in we can execute the same python file much faster. I have run my jupyter notebook in my local machine and also in IBM Watson Studio. Attaching the 2 times at the below table.

Local Machine Execution Time (seconds)	Cloud Execution Time(seconds)
1927.7977	44.85